Bachelor Thesis

My Bachelor Thesis

vorgelegt von

Marcel Pascal Stolin

32168

geboren am 03.04.1993 in Kamen

Verfasst am

Fraunhofer-Institut für Produktionstechnik und Automatisierung IPA

und

Hochschule der Medien Stuttgart

Erstprüfer: Prof. Walter Kriha
Zweitprüfer: Univ.-Prof. Dr.-Ing. habil. Marco Huber
Betreuung: M.Sc. Peter Trom

Stuttgart, October 27, 2020

Eidesstattliche Erklärung

Ich versichere hiermit, dass ich, Marcel Pascal Stolin, die vorliegende Bachelor Thesis selbstständig angefertigt, keine anderen als die angegebenen Hilfsmittel benutzt und sowohl wörtliche als auch sinngemäß entlehnte Stellen als solche kenntlich gemacht habe. Die Arbeit hat in gleicher oder ähnlicher Form noch keiner anderen Prüfungsbehörde vorgelegen.

Ort, Datum: _____

Unterschrift: _____
Marcel Pascal Stolin

Zusammenfassung

Hier kommt eine deutschsprachige Zusammenfassung hin.

Abstract

Abstract in English.

Contents

Notation

Konventionen

$$
\begin{array}{rl}
x & \text{Skalar} \\
\underline{x} & \text{Spaltenvektor} \\
\boldsymbol{x}, \underline{\boldsymbol{x}} & \text{Zufallsvariable/-vektor} \\
\hat{x}, \underline{\hat{x}} & \text{Mittelwert/-vektor} \\
x^*, \underline{x}^* & \text{Optimaler Wert/Vektor} \\
x_{0:k}, \underline{x}_{0:k} & \text{Folge von Werten } (x_0, x_1, \ldots, x_k) \text{ / Vektoren } (\underline{x}_0, \underline{x}_1, \ldots, \underline{x}_k) \\
\mathbf{A} & \text{Matrizen} \\
\mathcal{A} & \text{Mengen} \\
\preceq, \prec & \text{schwache/strenge Präferenzrelation} \\
\mathbb{R} & \text{Reelle Zahlen} \\
\mathbb{N} & \text{Natürliche Zahlen} \\
\blacksquare & \text{Ende eines Beispiels} \\
\square & \text{Ende eines Beweises}
\end{array}
$$

Operatoren

$$
\begin{array}{rl}
\mathbf{A}^{\mathrm{T}} & \text{Matrixtransposition} \\
\mathbf{A}^{-1} & \text{Matrixinversion} \\
|\mathbf{A}| & \text{Determinante einer Matrix} \\
|\mathcal{A}| & \text{Kardinalität der Menge } \mathcal{A} \\
\mathrm{pot}(\mathcal{A}) & \text{Potenzmenge von } \mathcal{A} \\
\mathrm{E}\{\cdot\} & \text{Erwartungswertoperator} \\
\mathcal{O}(g) & \text{O-Kalkül entsprechend der Landau-Notation bei welcher}
\end{array}
$$

beispielsweise $f(x) \in \mathcal{O}(g(x))$ besagt, dass die Funktion $f(x)$ die Komplexität $\mathcal{O}(g(x))$ besitzt

Spezielle Funktionen

$$
\begin{array}{rl}
\mathrm{Pr}(\mathcal{E}) & \text{Wahrscheinlichkeitsmaß, welches die Wahrscheinlichkeit}
\end{array}
$$

angibt, dass Ereignis $\mathcal{E}$ eintritt

$p(\underline{x})$ (Wahrscheinlichkeits-)Dichtefunktion für kontinuierliche $\underline{x}$

und Zähldichte für diskrete $\underline{x}$

$p(\underline{x}|\underline{y})$ Bedingte Dichtefunktion

$P(\underline{x})$ (Wahrscheinlichkeits-)Verteilungsfunktion

$\mathrm{erf}(x)$ Gauß'sche Fehlerfunktion

$\exp(x)$ Exponentialfunktion e^x

$\mathcal{N}(\underline{x}; \hat{\underline{x}}, \mathbf{C}_x)$ Gaußdichte, d. h. Dichtefunktion eines normalverteilten Zufallsvektors $\underline{\boldsymbol{x}}$ mit Mittelwertvektor $\hat{\underline{x}}$ und Kovarianzmatrix $\mathbf{C}_x$

Chapter 1

Introduction

Storing huge amounts of data has become inexpensive in recent years, but processing it, requires parallel computations on clusters with multiple machines [CZ18]. Complex computation operations rely on a IT infrastructure with the ability to perform operations on scale.

In order to achieve high scalability, computing systems need to adapt dynamically to demands and conditions of the workload.

1.1 Problem statement

ETL [1] operations are compute-intensive. During the execution of analytic applications, performance thresholds can be reached and the computing system can become out-of-order.

In addition, many algorithms profit from data-parallelism.

1.2 Research questions

1.3 Thesis structure

---

1   Extract transform load

Chapter 2

Related Work

Chapter 3

Background

In this chapter bla bla

3.1 Autonomic computing

Autonomic computing is an architecture for computing systems to enable the ability to manage themselves in accordance to high level objectives configured by administrators [KC03]. These computing systems dynamically adapt to demands and conditions of the workload [KC03]. An intelligent control-loop is responsible to collect all important details of the computing system and make decision according to the collected details. To automate the tasks, the intelligent control-loop is organized into four categories:

- **Self-configuring** - Components in the environment have to adapt dynamically to system changes using policies. For example, deploying or removing new components.

- **Self-healing** - If system errors have being detected, the control-loop has to perform policy-based actions without disrupting the environment.

- **Self-optimizing** - The control-loop has to monitor the resources and should adapt to changes dynamically.

- **Self-protecting** - Detection and protection against threats.

An autonomic computing environment consists of an autonomic manager, managed-resources and a knowledge-base.

3.1.1 Managed resources

Managed resources are software or hardware components in the computing environment. For example, a managed resource can be a database, service, application, server or different entity. Each managed resource implements

an interface to enable the autonomic manager to communicate with the managed resource.

These interface are called touchpoints.

3.1.2 Autonomic manager

The autonomic manager implements an intelligent control-loop to collect system metrics from the managed resources and acts according to the collected details. It can only make adjustments within it's own scope and uses policies to make decisions of what actions have to be executed to accommodate the objectives. To be self-managing, the autonomic manager has to implement the following four automated functions.

- **Monitor** - The monitor function is responsible to collect the needed metrics from all managed resources and applies aggregation and filter operations to the collected data. After that the function reports the metrics.

- **Analyze** - To determine if changes have to be made to the computing system, the collected data has to be analyzed.

- **Plan** - If changes have to be made, an appropriate change plan has to be generated. A change plan consists of actions that are needed to achieve the configured goals and objectives. The change plan needs to be forwarded to the execute function.

- **Execute** - The execute function applies all necessary changes to the computing system.

Multiple autonomic manager can exist in an autonomic computing environment to perform only certain parts. For example, there can be one autonomic manager which is responsible to monitor and analyze the system and another autonomic manager to plan and execute. To create a complete and closed control-loop, multiple autonomic manager can be composed together.

3.2 Docker

3.3 Apache Spark

Apache Spark is an open-source computing framework for parallel data processing on a large computer cluster. Spark manages the available resources and distributes computation tasks across a cluster to perform big-data processing operations at large scale [CZ18]. Before Spark was developed, Hadoop MapReduce [DG10] was the framework of choice for parallel operations on a computer cluster [ZCF+10]. Spark accomplished to outperform Hadoop by

10x for iterative Machine Learning [ZCF$^+$10]. It is implemented in Scala[1], a
JVM-based language and provides a programming interface for Scala, Java[2],
Python[3] and R[4]. In addition, Spark includes an interactive SQL shell and
libraries to implement Machine Learning and streaming applications [CZ18].
It was developed in 2009 as the Spark research project at UC Berkeley and
became an Apache Software Foundation project in 2013 [CZ18].

3.3.1 Spark programming model

Spark provides resilient distributed datasets (RDDs) as the main abstrac-
tation for parallel operations [ZCF$^+$10]. Core types of Spark's higher-level
structured API are built on top of RDDs [CZ18] and will automatically be
optimized by Spark's Catalyst optimizer to run operations quick and efficient
[Luu18].

TODO: Master-Slave Architektur + Bild

Resilient distributed datasets: Resilient distributed datasets are fault-
tolerant, parallel data structures to enable data sharing across cluster ap-
plications [ZCD$^+$12]. They allow to express different cluster programming
models like MapReduce, SQL and batched stream processing [ZCD$^+$12].
RDDs have been implemented in Spark and serve as the underlying data
structure for higher level APIs (Spark structured API) [ZCD$^+$12]. RDD's
are a immutable, partitioned collection of records and can only be initiated
through transformations (e.g. map, filter) on data or other RDD's. An
advantage of RDDs is, that they can be recovered through lineage. Lost
partitions of an RDD can be recomputed from other RDDs in parallel on
different nodes [ZCD$^+$12]. RDDs are lower level APIs and should only be
used in applications if custom data partitioning is needed [CZ18]. It is rec-
ommended to use Sparks structured API objects instead. Optimizations for
RDDs have to be implemented manually while Spark automatically optimize
the execution for structured API operations [CZ18].

Spark structured API: Spark provides high level structured APIs for
manipulating all kinds of data. The three distributed core types are Datasets,
DataFrames and SQL Tables and Views [CZ18]. Datasets and DataFrames
are immutable, lazy evaluated collections that provide execution plans for
operations [CZ18]. SQL Tables and Views work the same way as DataFrames,
except that SQL is used as the interface instead of using the DataFrame
programming interface [CZ18]. Datasets use JVM types and are therefore
only available for JVM based languages. DataFrames are Datasets of type
Row, which is the Spark internal optimized format for computations. This
has advantages over JVM types which comes with garbage collection and
object instantiation [CZ18].

---

1   Scala programming language. https://www.scala-lang.org/
2   Java programming language. https://www.oracle.com/java/
3   Python programming language. https://www.python.org/
4   R programming language. https://www.r-project.org/

Spark Catalyst: Spark also provides a query optimizer engine called Spark Catalyst. Figure 3.1 illustrates how the Spark Catalyst optimizer automatically optimizes Spark applications to run quickly and efficient. Before executing the user's code, the Catalyst optimizer translates the data-processing logic into a logical plan and optimizes the plan using heuristics [Luu18]. After that, the Catalyst optimizer converts the logical plan into a physical plan to create code that can be executed [Luu18].

Logical plans get created from a DataFrame or a SQL query. A logical plan represents the data-processing logic as a tree of operators and expressions where the Catalyst optimizer can apply sets of rule-based and cost-based optimizations [Luu18]. For example, the Catalyst can position a filter transformation in front of a join operation [Luu18].

From the logical plan, the Catalyst optimizer creates one ore more physical plans which consist of RDD operations [CZ18]. The cheapest physical will be generated into Java bytecode for execution across the cluster [Luu18].

Figure 3.1: Optimization process of the Spark Catalyst - Source: Authors own model, based on [Luu18].

3.3.2 Spark application architecture

Figure 3.2: Overview of a Spark cluster architecture - Source: Authors own model, based on [Apa20].

Figure 3.2 illustrates the main architecture of a Spark cluster. The architecture follows the master-worker model where the Spark driver is the master and the Spark executors are the worker [Luu18].

Spark driver: The Spark driver is a JVM process on a physical machine and responsible to maintain the execution of a Spark application [CZ18]. It

coordinates the application tasks onto each available executor [Luu18]. To get launch executors and get physical resources, the Spark driver interacts with the cluster manager [CZ18, Luu18].

Spark Executor: A Spark executor performs the tasks given by the Spark driver [CZ18]. It runs as a JVM process and runs until the Spark application finishes [Luu18]. After the executor finishes, it reports back to the Spark driver [CZ18]. Each task will be performed on a separate CPU core to enable parallel processing [Luu18].

Cluster manager: The cluster manager is an external service that orchestrates the work between the machines in the cluster [Luu18, Apa20]. The cluster manager knows about the resources of each worker and decides on which machine the Spark driver process and the executor processes run [Luu18, CZ18]. Spark supports different services that can run as the cluster manager: Standalone, Apache Mesos[5], Hadoop YARN[VMD+13] and Kubernetes[6] [Apa20]. The cluster manager provides three different deploy modes for acquiring resources in the cluster.

TODO: Explain standalone

- Cluster mode

- Client mode

- Local mode

To run an application in cluster mode, the user has to submit a precompiled JAR, python script or R script to the cluster manager [CZ18]. After that, the cluster manager starts the driver process and executor processes exclusively for the Spark application on machines inside the cluster [CZ18, Luu18].


Figure 3.3: Spark's cluster mode - Source: Authors own model, based on [CZ18].

The difference between the client mode and the cluster mode is that, the driver process runs on the client machine outside of the Spark cluster [CZ18].

The local mode starts a Spark application on a single computer [CZ18]. It is not recommended to use the local mode in production, instead it should be used for testing Spark applications or learning the Spark framework [CZ18].

---

5   Apache Mesos. https://mesos.apache.org/

6   Kubernetes. https://kubernetes.io/

Figure 3.4: Spark's client mode - Source: Authors own model, based on [CZ18].

3.3.3 Spark application implementation

The concept of a Spark application consists of calling transformations and actions. A transformation creates a DataFrame or a Dataset, the logical data structures of a Spark application. The computation of a Spark application gets processed when an action gets called in the application. The transformations of a Spark application build up a directed acyclic graph (DAG) of instructions. By calling an action, the DAG will break down into stages and tasks to create a single job for execution [CZ18].

```python
# Initialize a SparkSession
sparkSession = SparkSession\
    .builder\
    .getOrCreate()

# Create a dataframe with a transformation
dataframe = sparkSession.range(1, 1000)
# Apply another transformation
dataframe = dataframe.filter(dataframe.id % 2 == 0)
# Call an action
count = dataframe.count()
```

Listing 3.1: Example of a Python3 Spark application

Listing 3.1 demonstrates an example implementation of a Spark application. At first a SparkSession gets initialized. Each Spark application must include a SparkSession to initialize the application driver and executors [CZ18]. In Addition, the SparkSession provides an API for data-processing logic and configuration of the Spark application [Luu18]. After that, a DataFrame gets created with the range transformation to include each number from 1 to 1000 in the DataFrame. Next, a filter transformation is applied on the DataFrame to sort out any odd number. At the end, the number of rows gets saved in a variable with the count action.

The Spark framework provides a spark-submit executable to launch a Spark application inside a cluster.

<div style="float:left">**TODO: Create**
TODO: Tabelle
listing image
von Luu18 vll als
Anhang(Table 2-1
Subdirectories
Inside the
spark-2.1.1-bin-
hadoop2.7
Directory)</div>

```
$SPARK_HOME/bin/spark-submit \
    --master spark://spark-master:7077 \
    application.py
```

Listing 3.2: Execution of a Spark Python application using the spark-submit executable

Listing 3.2 provides an example how the spark-submit executable can be used to launch a Spark Python application.

3.3.4 Spark standalone cluster deployment

The standalone mode is a basic cluster-manager build specifically for Spark. It is developed to only run Spark but supports workloads at large scale [CZ18].

TODO: Why only standalone

Spark provides build-in scripts to start a master node and worker nodes in standalone mode. ABC demonstrates how a master node and worker node gets launched using the build-in scripts.

3.4 RAPIDS accelerator for Apache Spark

The RAPIDS accelerator for Apache Spark is a plugin suite to enable GPU acceleration for computing operations on Apache Spark 3.x [Spa20]. To accelerate computing operations, it uses the RAPIDS[7] libraries [Spa20] and extends the Spark programming model (see Subsection 3.3.1) [McD20].

3.4.1 Extension of the Spark programming model

The plugin suite extends the Spark programming model with a new DataFrame based on Apache Arrow[Fou20] data structures and the Catalyst optimizer to generate GPU-aware query plans [McD20].

Apache arrow is a data platform to build high performance applications that work with large dataset's and to improve analytic algorithms. A component of Apache Arrow is the Arrow Columnar Format, an in-memory data structure specification for efficient analytic operations on GPUs and CPUs [Fou20].

Spark's DataFrame and SQL use the RAPIDS APIs to run tranformations and actions on a GPU. The Spark Catalyst optimizer identifies operator in a query plan that are supported by the RAPIDS APIs. To execute the query plan, these operators can be scheduled on a GPU within the Spark cluster [McD20]. If operators are not supported by the RAPIDS APIs, a physical plan for CPUs will be generated by the Catalyst optimizer to execute RDD operations [McD20]. Figure 3.5 illustrates, how a query plan gets executed with the RAPIDS accelerator for Spark enabled.

3.4.2 GPU aware scheduling

---

7 Open GPU Data Science - https://rapids.ai/
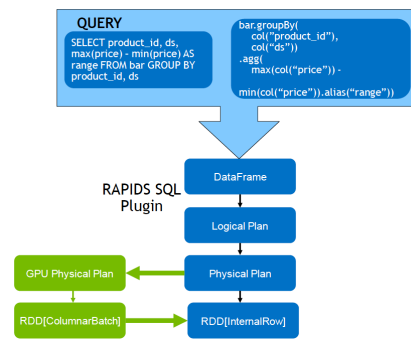
Figure 3.5: Query plan execution with RAPIDS accelerator for Apache Spark - Source: Authors own model, based on [McD20].

3.5 Prometheus

Prometheus is an open-source monitoring and alerting system [Pro20]. To collect and store data, Prometheus supports a multi-dimensional key-value pair based data model which can be analyzed in real-time using the PromQL query language [SP20]. It follows the pull-based approach to scrape metrics from hosts and services [BP19].

Figure 3.6: Pull-based approach to scrape metrics - Source: [BP19].

As Figure 3.6 demonstrates, a pull-based monitoring system scrapes metrics from services which makes them available for the monitoring system. In this case, the monitoring system needs a list of all hosts and services to monitor [BP19].
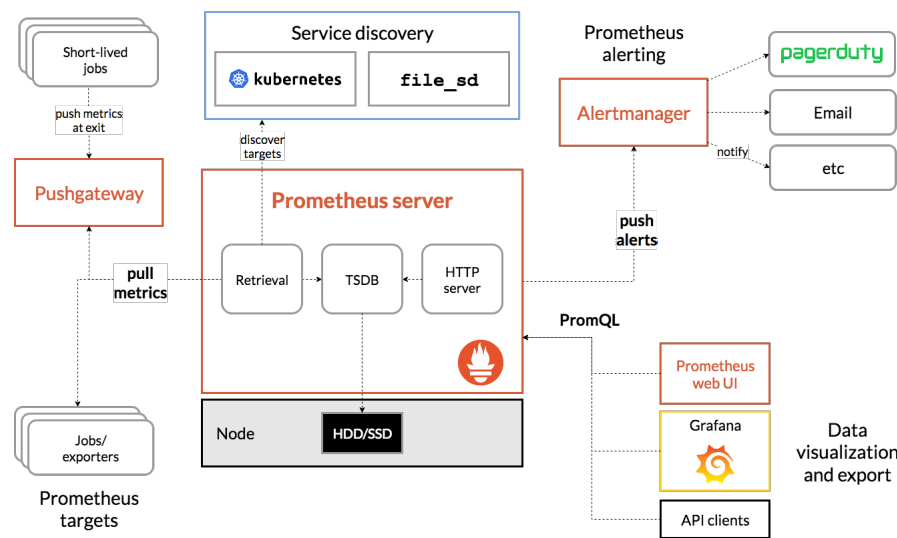
A monitoring system keeps track of the health status of components in a computing cluster. Because of the complexity of modern computing clusters, the effort is to high to observe components manually [BP19].

A metric is an observed property of software or hardware, e.g. the usage of a CPU core. To measure metrics, data-point's will be recorded at a fixed time interval. A data-point is a pair of a value and a timestamp. The combination of multiple data-point's is called a time-series [Tur16].

3.5.1 Prometheus architecture

TODO: Bilder für components

Figure 3.7 illustrates the high-level architecture of Prometheus. The Prometheus ecosystem provides multiple components. Components can be optional, depending on the monitoring needs of the environment [BP19]. The main components of a Prometheus system are Prometheus server, Alertmanager, service discovery, exporters, Pushgateway and visualization tools [Pro20].

Figure 3.7: Prometheus high-level architecture - Source: Authors own model, based on [Pro20, BP19].

Prometheus server: The Prometheus server is the main component of a Prometheus system. It is responsible to collect metrics as time-series data from targets and stores the collected data in the built-in TSDB [BP19]. Prometheus uses the concept of scraping to collect metrics from a target. A target host has to expose an endpoint to make metrics available in the Prometheus data format [SP20]. Additionally, the Prometheus server triggers alerts to the Alertmanager if a configured condition becomes true [Pro20].

Alertmanager: If an alerting rule becomes true, the Prometheus server generates an alert and pushes it to the Alertmanager. The Alertmanager generates notifications from the received alerts. A notification can take multiple forms like emails or chat messages. Webhooks can be implemented to trigger custom notifications [BP19].

Service discovery: As mentioned before, Prometheus follows a pull-based approach to fetch metrics from a target. To know about all targets, Prometheus needs a list of the corresponding hosts. The service discovery manages the complexity of maintaining a list of hosts manually in an changing infrastructure [BP19].

Exporters: If an application does not support an endpoint for Prometheus, an exporter can be used to fetch metrics and make them available to the Prometheus server. An exporter is a monitoring agent running on a target host that fetches metric from the host and exports them to the Prometheus server [SP20].

Pushgateway: If a target is not designed to be scraped, metrics can be pushed against the Pushgateway[Pro20]. The Pushgateway converts the

data into the Prometheus data format and passes them to the Prometheus server [SP20].

Visualization: Prometheus supports various tools for virtualization of the scraped data. Grafana[8] is one of the widely used tools for this occasion.

3.5.2 Monitoring Docker container

TODO: Wahrscheinlich erst nach der implementation

3.6 Gitlab CI/CD

3.7 K-MEANS

3.8 Naive Bayes Classifier

3.9 Scaling heat

3.10 KHP

---

8   Grafana: The open observability platform - https://grafana.com/

Chapter 4

Spark cluster implementation

4.1 Cluster architecture

Bibliography

[Apa20]   *Spark 3.0.1 Documentation.* `https://spark.apache.org/docs/3.0.1/`. Version: Oktober 2020

[BP19]   Bastos, Joel ; Pedro, Araujo: *Hands-On Infrastructure Monitoring with Prometheus.* Packt Publishing, 2019. – ISBN 9781789612349

[CZ18]   Chambers, Bill ; Zaharia, Matei: *Spark: The Definitive Guide.* 1st. O'Reilly Media, 2018. – ISBN 9781491912218

[DG10]   Dean, Jeffrey ; Ghemawat, Sanjay: MapReduce: A Flexible Data Processing Tool. In: *Commun. ACM* 53 (2010), Januar, Nr. 1, 72–77. `http://dx.doi.org/10.1145/1629175.1629198`. – DOI 10.1145/1629175.1629198. – ISSN 0001–0782

[Fou20]   Foundation, The A.: *Arrow. A cross-language development platform for in-memory data.* `https://arrow.apache.org/`. Version: Oktober 2020

[KC03]   Kephart, J. O. ; Chess, D. M.: The vision of autonomic computing. In: *Computer* 36 (2003), Nr. 1, S. 41–50

[Luu18]   Luu, Hien: *Beginning Apache Spark 2: With Resilient Distributed Datasets, Spark SQL, Structured Streaming and Spark Machine Learning library.* 1st. Apress, 2018. – ISBN 9781484235799

[McD20]   McDonald, Carol: *ACCELERATING APACHE SPARK 3.X.* 2020

[Pro20]   *Prometheus Online Documentation.* `https://prometheus.io/docs/`. Version: Oktober 2020

[SP20]   Sabharwal, Navin ; Pandey, Piyush: *Monitoring Microservices and Containerized Applications: Deployment, Configuration, and Best Practices for Prometheus and Alert Manager.* Apress, 2020. – ISBN 9781484262160

[Spa20] *Spark-Rapids Online Documentation.* `https://nvidia.github.io/spark-rapids/`. Version: Oktober 2020

[Tur16] TURNBULL, James: *The Art of Monitoring.* Turnbull Press, 2016

[VMD+13] VAVILAPALLI, Vinod K. ; MURTHY, Arun C. ; DOUGLAS, Chris ; AGARWAL, Sharad ; KONAR, Mahadev ; EVANS, Robert ; GRAVES, Thomas ; LOWE, Jason ; SHAH, Hitesh ; SETH, Siddharth ; SAHA, Bikas ; CURINO, Carlo ; O'MALLEY, Owen ; RADIA, Sanjay ; REED, Benjamin ; BALDESCHWIELER, Eric: Apache Hadoop YARN: Yet Another Resource Negotiator. New York, NY, USA : Association for Computing Machinery, 2013 (SOCC '13). – ISBN 9781450324281

[ZCD+12] ZAHARIA, Matei ; CHOWDHURY, Mosharaf ; DAS, Tathagata ; DAVE, Ankur ; MA, Justin ; MCCAULY, Murphy ; FRANKLIN, Michael J. ; SHENKER, Scott ; STOICA, Ion: Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing. In: *9th USENIX Symposium on Networked Systems Design and Implementation (NSDI 12).* San Jose, CA : USENIX Association, April 2012. – ISBN 978–931971–92–8, 15–28

[ZCF+10] ZAHARIA, Matei ; CHOWDHURY, Mosharaf ; FRANKLIN, Michael J. ; SHENKER, Scott ; STOICA, Ion: Spark: Cluster Computing with Working Sets. In: *Proceedings of the 2nd USENIX Conference on Hot Topics in Cloud Computing.* USA : USENIX Association, 2010 (HotCloud'10), S. 10

List of Figures